64-bit Processors: The Future of Computing

Introduction

A 64-bit processor refers to a microprocessor that can process data and instructions in chunks of 64 bits. Microprocessors that can handle 64 bits perform a larger number of calculations per second compared to 32-bit processors. Typical variations of the term include 64-bit CPU, 64-bit computing and 64-bit microprocessor. 64-bit processors are the latest and most advanced technology in the world of computing. [1]

The 64-bit computer originated in 1961 when IBM created the IBM 7030 Stretch supercomputer. However, it was not put into use in home computers until the early 2000s. Microsoft released a 64-bit version of Windows XP to be used on computers with a 64-bit processor. Windows Vista, Windows 7, and Windows 8 also come in 64-bit versions.

On a computer with a 64-bit processor, you cannot run a 16-bit legacy program. Many 32-bit programs works with a 64-bit processor and operating system, but some older 32-bit programs may not function properly, or at all, due to limited or no compatibility. [2]

Every microprocessor handles two key criteria: memory addresses and the corresponding data or instruction present at each memory address. Thus, microprocessors are tied to memory where applications and data are executed, and the number of bits available on a processor define the memory size and space, as well as the volume of data that a processor can handle per unit of time. This has become a de facto way of classifying or categorizing modern processors.

A 64-bit processor uses internal registers -- temporary storage locations within the processor -- that are 64 bits wide. This typically corresponds to an address bus and data bus that are also 64 bits wide. The address bus is the pathway of electrical signals used to determine the device or memory address that the processor is attempting to access. The data bus is the pathway used to exchange data with the intended address. Processors also include signaling for a third control bus, but this bus is typically a unique collection of discrete, or individual, digital signals and does not operate like an address or data bus.
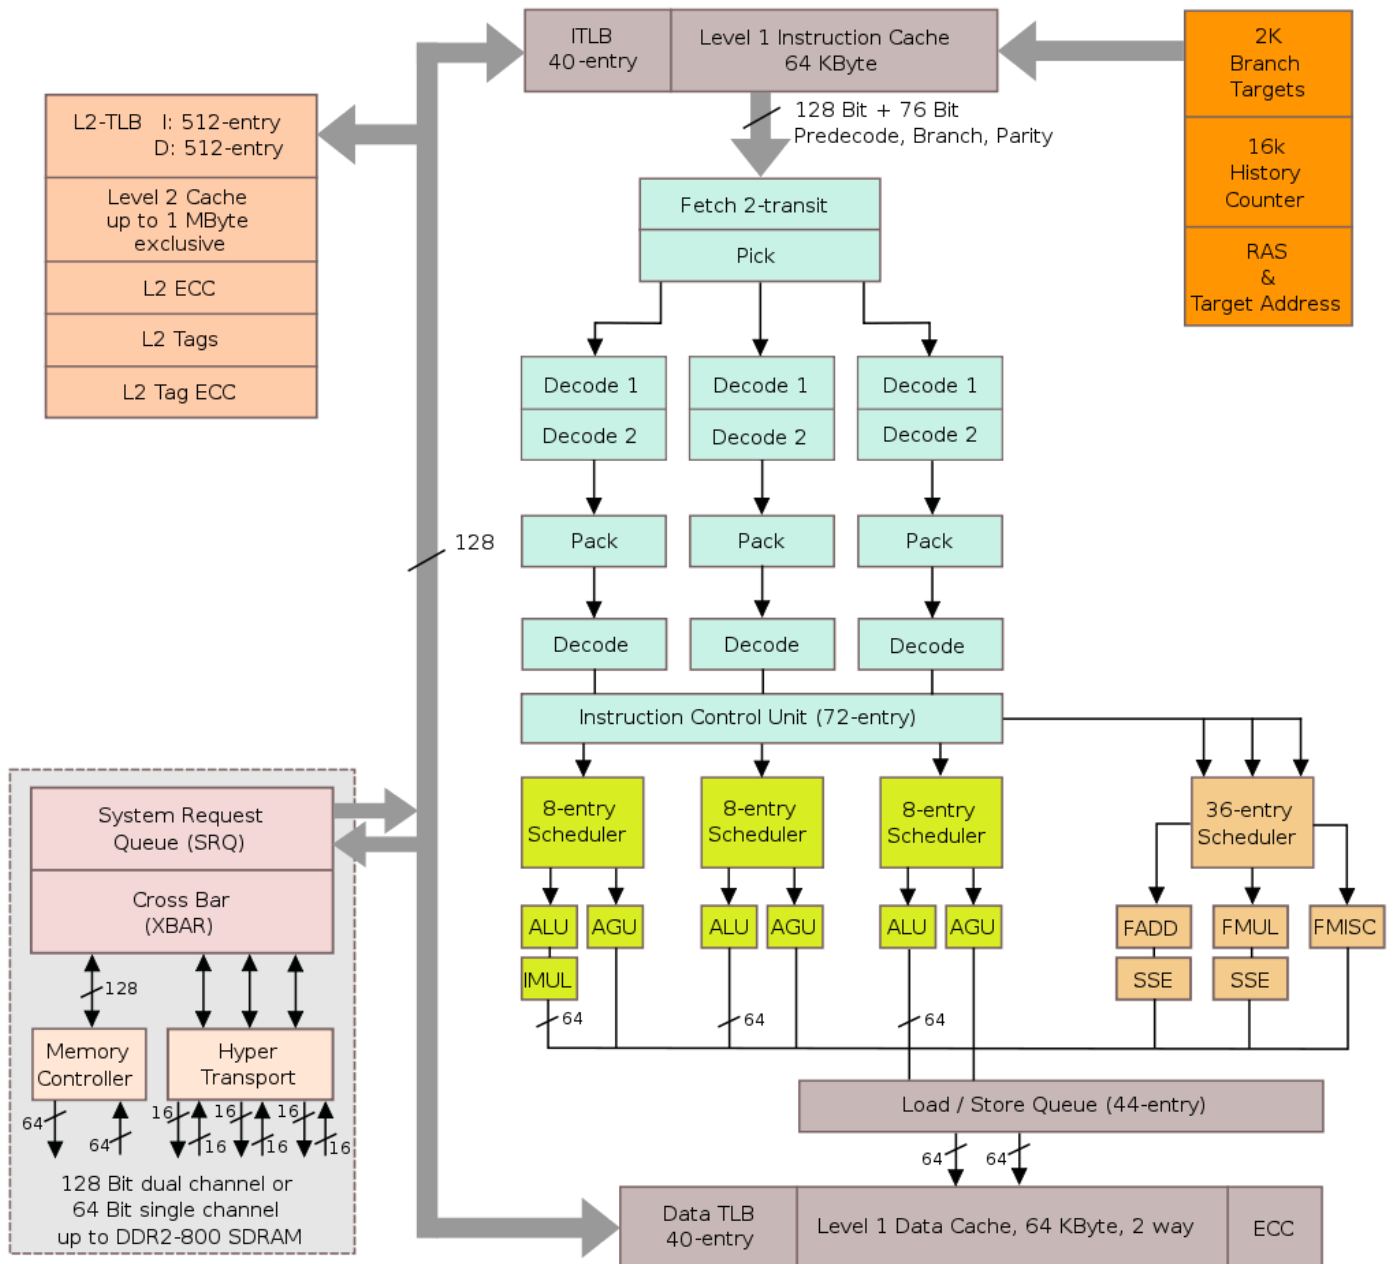
Internal Architecture

AMD K8 Architecture

Fig: Internal Architecture of a Athlon 64 processor [3]

The internal architecture of a 64-bit processor is one of the most important factors that determine its performance and capabilities. Compared to 32-bit processors, 64-bit processors have a larger data bus and address bus, which allows them to handle larger data and address more memory. In addition, there are more registers inside the processor that store data for processing, which is a small and fast memory location.

The wider data bus and address bus of a 64-bit processor allows more data to be transferred per clock cycle than a 32-bit processor. This means they can do more at the same time and do it faster. For example, a 32-bit processor can handle data in 32-bit chunks, while a 64-bit processor can handle data in 64-bit chunks. This allows 64-bit processors to process large amounts of data simultaneously, making them ideal for demanding applications and multitasking.

In addition, 64-bit processors have wider data buses and address buses, allowing them to access more memory than 32-bit processors. A 32-bit processor can access up to 4 GB of RAM, while a 64-bit processor can access up to 16 GB of RAM. This means that 64-bit processors can handle larger amounts of data and run more resource-intensive applications without experiencing significant slowdowns or crashes. This is especially useful for professionals who work with large amounts of data, such as data scientists and engineers.

The number of registers on a 64-bit processor is a significant advantage over a 32-bit processor. The computer stores the data currently being processed by the processor, allowing it to perform calculations and execute instructions. 64-bit processors have more registers than 32-bit processors, which means they can handle more complex instructions and perform more calculations at the same time. This results in faster processing and better performance.

In addition to the advantages mentioned above, the wide data bus and address bus of the 64-bit processor allows the use of a wider instruction set than the 32-bit processor. An instruction set is a set of instructions that a processor can execute. With a wider instruction set, a 64-bit processor can handle more complex instructions and perform more efficient data processing.

Another important aspect of the internal architecture of 64-bit processors is the ability to manage virtual memory. Virtual memory is a way to allow a computer to access more physical memory. This is useful for running large applications that require more memory than the physical system. A 64-bit processor can manage virtual memory more efficiently than a 32-bit processor, which means it can run more programs simultaneously and handle more data-intensive applications.

The internal architecture of a 64-bit processor can vary depending on the specific design and manufacturer, but here are the general overview of the key components and features typically found in such processors.

1. **Instruction Fetch and Decode Unit**: This unit is responsible for fetching instructions from memory and decoding them into a format that the processor can understand and execute
   - Program Counter (PC): Specifies the memory address of the next instruction to be fetched (64 bits).
   - Opcode: Specifies the operation to be performed (8 to 16 bits).
   - Source Register 1: Identifies the first source register (5 to 6 bits).
   - Source Register 2: Identifies the second source register (5 to 6 bits).
   - Destination Register: Identifies the destination register (5 to 6 bits).
   - Immediate Value: Specifies an immediate constant or operand (16 to 32 bits).
2. **Control Unit(CU)**: The control unit manages the flow of data and instructions within the processor. It coordinates various components to ensure proper execution of instructions.
   - Control Signals: Various control signals that manage the flow of data and instructions within the processor.
3. **Arithmetic Logic Unit (ALU)**: The ALU performs arithmetic and logical operations, such as addition, subtraction, multiplication, division, bitwise operations, and comparisons. It operates on 64-bit data or multiple smaller data units in parallel.
   - Opcode: Specifies the arithmetic or logical operation to be performed (4 to 8 bits).
   - Operand 2: Specifies the second operand (64 bits).
   - Result: Stores the result of the operation (64 bits).
   - Operand 1: Specifies the first operand (64 bits).
4. **Register File**: A register file is a small, high-speed memory within the processor that stores frequently accessed data. It consists of multiple registers, typically 64 bits wide, which can hold operands, intermediate results, and addresses.
   - Register Address: Identifies a specific register (5 to 6 bits).
   - Data: Holds the value stored in the register (64 bits).
5. **Floating-Point Unit (FPU)**: The FPU handles floating-point arithmetic operations, including addition, subtraction, multiplication, division, and transcendental functions. It supports 64-bit floating-point numbers (often using the IEEE 754 standard) and can have its own set of registers.

- Opcode: Specifies the floating-point operation to be performed (8 to 16 bits).
- Floating-Point Register 1: Identifies the first floating-point register (5 to 6 bits).
- Floating-Point Register 2: Identifies the second floating-point register (5 to 6 bits).
- Floating-Point Result Register: Identifies the register to store the result (5 to 6 bits).

6. **Cache Memory**: Processors typically have multiple levels of cache memory, including L1, L2, and sometimes L3 caches. These caches store frequently accessed instructions and data, providing faster access than main memory
    - Cache Line Address: Specifies the memory address or index of the cache line (64 bits).
    - Data: Stores the data retrieved from or written to the cache (64 bits).

7. **Memory Management Unit (MMU)**: The MMU translates virtual memory addresses used by programs into physical memory addresses. It helps manage memory protection, virtual memory, and cache coherence.
    - Virtual Address: Specifies the virtual memory address (64 bits).
    - Physical Address: Stores the corresponding physical memory address (64 bits).

8. **Pipelining**: Processors often employ pipelining techniques to increase instruction throughput. Pipelining breaks down instruction execution into multiple stages, allowing multiple instructions to be processed simultaneously. These stages can include instruction fetch, decode, execution, memory access, and write-back.
    - Pipeline Registers: Store intermediate results and data between pipeline stages (64 bits).

9. **Branch Prediction Unit**: The branch prediction unit analyzes branch instructions to predict the outcome, improving the efficiency of instruction execution by reducing pipeline stalls caused by branch instructions.
    - Branch Address: Specifies the target address of a branch instruction (64 bits).
    - Prediction Result: Indicates the predicted outcome of the branch (1 bit).

10. **Out-of-Order Execution**: Some 64-bit processors utilize out-of-order execution, which allows instructions to be executed in a different order than

they appear in the program. This technique leverages available resources and reduces the impact of dependencies, improving overall performance.

- Opcode: Specifies the operation to be performed.
- Source/Destination Registers: Identifies the source and destination registers for the operation.
- Immediate Values: Provides immediate data for the instruction.
- Control Flow Information: Contains information related to branch prediction and control flow.

11. **Bus Interface**: The processor has interfaces to communicate with other system components, such as main memory, peripheral devices, and other processors in a multi-processor system. These interfaces can include address buses, data buses, and control buses.

- Address Bus: Transfers memory or device addresses (64 bits).
- Data Bus: Transfers data between the processor and memory or devices (64 bits).
- Control Bus: Carries control signals for data transfer and synchronization.

12. **Vector Processing Unit**: Some 64-bit processors feature a vector processing unit, also known as SIMD (Single Instruction, Multiple Data), which can perform parallel operations on multiple data elements simultaneously. This unit is optimized for tasks that involve large amounts of data processing, such as multimedia and scientific applications.

- Vector Operation: Specifies the operation to be performed on the vectors.
- Vector Length: Specifies the length or size of the vector operands.
- Source/Destination Vector Registers: Identifies the source and destination vector registers.
- Control Information: Contains control flags or other relevant information for vector processing.

13. **Memory Controller**: The memory controller handles the communication between the processor and the main memory. It manages the read and write operations, memory access timings, and data transfer between the processor and memory modules.

14. **On-Chip Cache Coherence Protocol**: In multi-processor systems, where multiple processors share a common memory, an on-chip cache coherence

protocol ensures that all processors observe a consistent view of memory. It manages data coherency across multiple processor caches to prevent conflicts and maintain data integrity.

15. **Power Management Unit**: Modern processors often include a power management unit responsible for regulating power consumption and optimizing energy efficiency. It dynamically adjusts clock frequencies, voltage levels, and puts parts of the processor into low-power states when not in use.

16. **Branch Target Buffer**: The branch target buffer is a small, high-speed cache that stores the predicted target addresses of branch instructions. It helps improve branch prediction accuracy by storing historical branch behavior and reducing pipeline stalls caused by incorrect branch predictions.

In summary, the internal architecture of a 64-bit processor plays an important role in its performance and capabilities. A wider data bus and address bus, a larger number of registers, greater memory access, and the ability to manage virtual memory allow them to handle more data and perform more complex tasks than their 32-bit counterparts. This has made 64-bit processors the industry standard and will continue to grow and improve in the future.

Comparision between 32-bit and 64-bit processor

The difference between 32-bit and 64-bit processors are: [4][5]

Parameter	32-bit processors	64-bit processors
Addressable space	It has 4 GB addressable space	64-bit processors have 16 exabytes addressable space
Application support	64-bit applications and programs won't work	32-bit applications and programs will work
OS support	Need a 32-bit operating system.	It can run on 32 and the 64-bit operating system.
Support for multi-tasking	Not an ideal option for stress testing and multi-tasking.	Works best for performing multi-tasking and stress testing.
OS and CPU requirement	32-bit operating systems and applications require 32-bit CPUs	64-bit OS demands 64-bit CPU, and 64-bit applications require 64-bit OS and CPU.
System available	Support Windows 7, 8 Vista, XP, and, Linux.	Windows XP Professional, Windows Vista, Windows 7, Windows 8,Windows 10, Linux, and Mac OS X.
Memory limits	32-bit systems limited to 3.2 GB of RAM 32 bit Windows. It addresses limitation doesn't allow you to use full 4GB of Physical memory space.	64-bit systems will enable you to store up to 17 Billion GB of RAM.
Performance	Limited by the maximum amount of RAM it can access	Can take advantage of more memory, enabling faster performance
Compatibility	Can run 32-bit and 16-bit applications	Can run 32-bit and 64-bit applications

Parameter	32-bit processors	64-bit processors
Hardware support	May not support newer hardware	Supports newer hardware with 64-bit drivers
Security	Limited security features	More advanced security features, such as hardware-level protection
Application support	Limited support for new software	Supports newer software designed for 64-bit architecture
Price	Less expensive than 64-bit OS	More expensive than 32-bit OS
Gaming	Can run high graphical games, but may not be as efficient as with 64-bit OS	Can run high graphical games and handle complex software more efficiently
Virtualization	Limited support for virtualization	Better support for virtualization

Advantages

There are 3 most obvious advantages of 64-bit processors over their 32-bit counterparts: extended address space, capacity increase, and larger number of general-purpose registers.[6][7][8][9]

The extended 64-bit address space theoretically allows the processor to work with 16 Ebytes (2^64) of physical memory within the flat memory model. Although contemporary 64-bit processors can address only 1 Tbyte (2^40) of memory in practice, it is still much higher than in the case of 32-bit addressing. A larger amount of available memory, in turn allows you to eliminate or greatly reduce sluggish operations of data swapping to the disk.

The larger number of registers, and increase of their size, allow the processor to handle large memory areas simultaneously, to handle variables and arrays more effectively, and to pass function arguments in registers instead of the stack.

Keep in mind that you must port your program using the 64-bit version of the compiler, and consider changes of the data model (new type capacities) to get a real performance gain on a 64-bit processor. Launching an application not adapted for the 64-bit platform might on the contrary cause great performance loss, depending on the specifics of the architecture of the processor being used.

But the greatest performance gain from moving to the 64-bit platform awaits those applications which handle large data arrays - these are CAD/CAE-like programs, data managers, programs to work with digital media and applied scientific applications. The performance gain might be of several hundred percent for this kind of software.

Thus, for instance, porting the service 'Oracle on Demand' by the Oracle company from the 32-bit platform to the 64-bit one, allowed reduction of the server response time (350 users working simultaneously) from 13.60 seconds to 2.35 seconds. When FASTCOM integrated the 64-bit platform based on Itanium 2 into their IT-infrastructure, they managed to eliminate issues caused by increased load on billing systems, and get a two-fold performance gain.

Limitations

Some of the limitations of 64-bit processors are: [10][11][12][13]

- **Industry catch-up**: Today's 64-bit processors cannot access all of the potential memory space that 64 bits offer.
- **Code Size**: Transitioning from 32-bit to 64-bit architectures often results in an increase in code size. This is primarily due to larger pointers and expanded data types. The larger code size can consume more memory and cache space, potentially impacting cache utilization and instruction fetch rates.
- **More memory occupancy for small tasks**: A 64-bit operating system computing takes up more memory space in comparison to 32-bit systems for the same data.
- **Memory Bandwidth**: With larger data paths and increased memory addressability, 64-bit processors may require higher memory bandwidth to efficiently handle data transfers. If the memory subsystem cannot provide sufficient bandwidth, it can become a bottleneck and limit the overall system performance.
- **Low availability of 64-bit drivers for older systems**: It is hard to find 64-bit drivers for older systems as there is a very less number of 64-bit drivers available for older hardware and systems.
- **Backward Compatibility**: While 64-bit processors can execute 32-bit code, running legacy 16-bit or 8-bit applications can pose challenges. These older software applications may require significant modifications or emulation layers to run on 64-bit processors, impacting compatibility and potentially introducing performance overhead.
- **Operating system limitations**: Some operating systems impose limits on the address space or the physical memory that can be used by 64-bit processors. For example, Windows Vista allows only 16 TB of address space and 4 PiB of physical memory.
- **Software compatibility issues**: Some software applications may not work well or at all on 64-bit systems, especially if they rely on 32-bit drivers or libraries. Some software may also have bugs or performance issues when running on 64-bit systems.

- **Cost and power consumption**: Generally, 64-bit processors are more expensive and consume more power than 32-bit processors, which may not be desirable for some users or applications

- **Security risks**: Some security features, such as address space layout randomization (ASLR), may be less effective on 64-bit systems due to the larger address space. Some malware may also exploit the increased memory space to hide or evade detection

- **Hardware compatibility issues**: Some hardware devices may not support 64-bit processors or may require firmware updates or adapters to work properly. Some hardware may also have performance degradation or reduced functionality on 64-bit systems

- **Lack of standardization**: There are different types of 64-bit architectures, such as x86-64, ARMv8, PowerPC, etc., which may have different features, instruction sets, and conventions. This may cause confusion or inconsistency among developers and users.

- **Memory Overhead**: 64-bit processors require more memory to store 64-bit addresses compared to 32-bit processors. This can lead to increased memory usage and potentially impact system performance, especially in memory-constrained environments or applications that do not require the full 64-bit address space.

- **Transition and Compatibility Challenges**: The transition from a 32-bit ecosystem to a 64-bit one can present challenges in terms of software compatibility, driver availability, and third-party support. Ensuring a smooth transition and compatibility across all software and hardware components can require effort and coordination from both software developers and hardware manufacturers.

Conclusion

In this report, we have discussed the concept and features of 64-bit processors, which are microprocessors that can process data and instructions in chunks of 64 bits. We have also compared 64-bit processors with 32-bit processors, which are limited to 4 GB of addressable memory space and perform fewer calculations per second. We have examined the internal architecture of 64-bit processors, which use 64-bit registers, address buses, and data buses to handle larger amounts of data and memory. We have also explored some of the limitations and challenges of 64-bit computing, such as compatibility issues, memory consumption, and security risks.

We have concluded that 64-bit processors offer significant advantages over 32-bit processors in terms of performance, memory capacity, and scalability. 64-bit processors have been used in supercomputers since the 1970s and in workstations and servers since the early 1990s. They have also become the norm for personal computers since the early 2000s. However, 64-bit computing is not without its drawbacks and limitations, and it requires careful design and implementation to ensure optimal results. Therefore, we recommend that users and developers should be aware of the benefits and challenges of 64-bit computing and choose the appropriate processor and software for their needs and goals.

References

1. TechTarget. (n.d.). What is a 64-Bit Processor (64-Bit Computing)? Retrieved from *https://www.techtarget.com/searchdatacenter/definition/64-bit-processor*↵

2. Computer Hope. (2021, October 7). What is the difference between a 32-bit and 64-bit system? Retrieved from *https://www.computerhope.com/issues/ch001498.htm#64bit*↵

3. Wikipedia. (2021, December 13). Athlon 64. Retrieved from *https://en.wikipedia.org/wiki/Athlon_64*↵

4. GeeksforGeeks. (2020, August 28). Difference between 32-bit and 64-bit operating systems. Retrieved from *https://www.geeksforgeeks.org/difference-32-bit-64-bit-operating-systems/*↵

5. Guru99. (2021, November 23). What is the difference between 32 bit vs 64 bit Operating Systems? Retrieved from *https://www.guru99.com/32-bit-vs-64-bit-operating-systems.html*↵

6. Ryzhkov, E., & Karpov, A. (2009). A 64-bit horse that can count. Retrieved from *https://pvs-studio.com/en/blog/posts/a0029*↵

7. Stackoverflow.com. (2010, March 26). 32 vs 64 bit - or what exactly does "processing information" mean? Retrieved from *https://stackoverflow.com/questions/2532129/32-vs-64-bit-or-what-exactly-does-processing-information-mean*↵

8. Karpov, A., & Ryzhkov, E. (2009). Lessons on development of 64-bit C/C++ applications. Retrieved from *https://pvs-studio.com/en/blog/lessons/*↵

9. Karpov, A. (2009). 20 issues of porting C++ code on the 64-bit platform. Retrieved from *https://pvs-studio.com/en/blog/posts/cpp/a0030/*↵

10. Scaler. (n.d.). 64-bit Operating System. Retrieved from *https://www.scaler.com/topics/64-bit-operating-system/*↵

11. Wikipedia. (2021, December 13). 64-bit computing. Retrieved from *https://en.wikipedia.org/wiki/64-bit_computing/*↵

12. Super User. (2009, November 3). What is the maximum amount of ram a 64bit machine can theoretically address? Retrieved from *https://superuser.com/questions/65191/what-is-the-maximum-amount-of-ram-a-64bit-machine-can-theoretically-address*↵

13. Super User. (2013, February 19). Memory limits in 16, 32 and 64 bit systems. Retrieved from *https://superuser.com/questions/556008/memory-limits-in-16-32-and-64-bit-systems*↵