
CHAPTER 3: *General-Purpose Processors: Software*



- 3.1 Introduction
- 3.2 Basic Architecture
- 3.3 Operation
- 3.4 Programmer's View
- 3.5 Development Environment
- 3.6 Application-Specific Instruction-Set Processors (ASIPs)
- 3.7 Selecting a Microprocessor
- 3.8 General-Purpose Processor Design
- 3.9 Summary
- 3.10 References and Further Reading
- 3.11 Exercises

3.1 Introduction

A general-purpose processor is a programmable digital system intended to solve computation problems in a large variety of applications. Copies of the same processor may solve computation problems in applications as diverse as communication, automotive, and industrial embedded systems. An embedded-system designer choosing to use a general-purpose processor to implement part of a system's functionality may achieve several benefits.

First, the unit cost of the processor may be very low, often a few dollars or less. One reason for this low cost is that the processor manufacturer can spread its NRE cost for the processor's design over large numbers of units, often numbering in the millions or billions. For example, Motorola sold nearly half a billion 68HC05 microcontrollers in 1996 alone (source: Motorola 1996 Annual Report).

Second, because the processor manufacturer can spread NRE cost over large numbers of units, the manufacturer can afford to invest large NRE cost into the processor's design, without significantly increasing the unit cost. The processor manufacturer may thus use

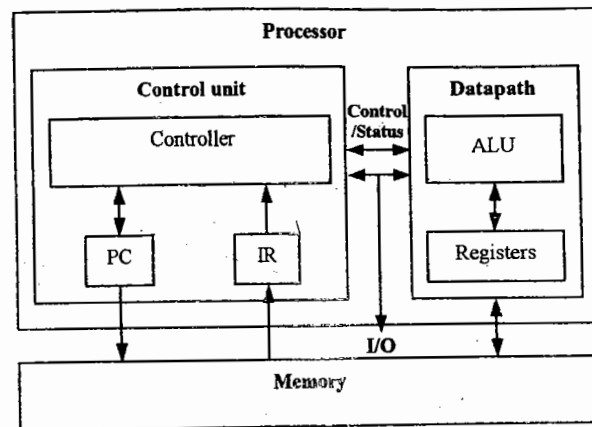


Figure 3.1: General-purpose processor basic architecture.

experienced computer architects who incorporate advanced architectural features, and may use leading-edge optimization techniques, state-of-the-art IC technology, and handcrafted VLSI layouts for critical components. These factors can improve design metrics like performance, size and power.

Third, the embedded system designer may incur low NRE cost, since the designer need only write software, and then apply a compiler and/or an assembler, both of which are mature and low-cost design technologies. Likewise, time-to-prototype and time-to-market will be short, since processor ICs can be purchased and then programmed in the designer's own lab. Flexibility will be great, since the designer can perform software rewrites in a straightforward manner.

3.2 Basic Architecture

A general-purpose processor, sometimes called a CPU (central processing unit) or a microprocessor, consists of a datapath and a control unit, tightly linked with a memory. We now discuss these components briefly. Figure 3.1 illustrates the basic architecture.

Datapath

The datapath consists of the circuitry for transforming data and for storing temporary data. The datapath contains an arithmetic-logic unit (ALU) capable of transforming data through operations such as addition, subtraction, logical AND, logical OR, inverting, and shifting. The ALU also generates status signals, often stored in a status register (not shown), indicating particular data conditions. Such conditions include indicating whether data is zero or whether an addition of two data items generates a carry. The datapath also contains registers capable

of storing temporary data. Temporary data may include data brought in from memory but not yet sent through the ALU, data coming from the ALU that will be needed for later ALU operations or will be sent back to memory, and data that must be moved from one memory location to another. The internal data bus carries data within the datapath, while the external data bus carries data to and from the data memory.

We typically distinguish processors by their size, and we usually measure size as the bit-width of the datapath components. A *bit*, which stands for binary digit, is the processor's basic data unit, representing either a 0 (low or false) or a 1 (high or true), while we refer to 8 bits as a *byte*. An *N-bit* processor may have *N*-bit-wide registers, an *N*-bit-wide ALU, an *N*-bit-wide internal bus over which data moves among datapath components, and an *N*-bit wide external-bus over which data is brought in and out of the datapath. Common processor sizes include 4-bit, 8-bit, 16-bit, 32-bit, and 64-bit. However, in some cases, a particular processor may have different sizes among its registers, ALU, internal bus, or external bus, so the processor-size definition is not an exact one. For example, a processor may have a 16-bit internal bus, ALU and registers, but only an 8-bit external bus to reduce pins on the processor's IC.

Control Unit

The control unit consists of circuitry for retrieving program instructions and for moving data to, from, and through the datapath according to those instructions. The control unit has a program counter (PC) that holds the address in memory of the next program instruction to fetch, and an instruction register (IR) to hold the fetched instruction. The control unit also has a controller, consisting of a state register plus next-state and control logic, as we saw in Chapter 2. This controller sequences through the states and generates the control signals necessary to read instructions into the IR, and control the flow of data in the datapath. Such flows may include inputting two particular registers into the ALU, storing ALU results into a particular register, or moving data between memory and a register. The controller also determines the next value of the PC. For a nonbranch instruction, the controller increments the PC. For a branch instruction, the controller looks at the datapath status signals and the IR to determine the appropriate next address.

The PC's bit-width represents the processor's address size. The address size is independent of the data word size; the address size is often larger. The address size determines the number of directly accessible memory locations, referred to as the *address space* or *memory space*. If the address size is M , then the address space is 2^M . Thus, a processor with a 16-bit PC can directly address $2^{16} = 65,536$ memory locations. We would typically refer to this address space as 64K, although if 1K = 1,000, this number would represent 64,000, not the actual 65,536. Thus, in computer-speak, 1K = 1,024.

For each instruction, the controller typically sequences through several stages, such as fetching the instruction from memory, decoding it, fetching operands, executing the instruction in the datapath, and storing results. Each stage may consist of one or more clock cycles. A clock cycle is usually the longest time required for data to travel from one register to another. The path through the datapath or controller that results in this longest time (e.g., from a datapath register through the ALU and back to a datapath register) is called the *critical*




Figure 3.2: Two memory architectures: (a) Harvard, (b) Princeton.

path. The inverse of the clock cycle is the clock frequency, measured in cycles per second, or Hertz (Hz). For example, a clock cycle of 10 nanoseconds corresponds to a frequency of $1/10 \times 10^{-9}$ Hz, or 100 MHz. The shorter the critical path, the higher the clock frequency. We often use clock frequency as a means of comparing processors, especially different versions of the same processor, with higher clock frequency implying faster program execution. However, using clock frequency is not always an accurate method for comparing processor speeds.

Memory

While registers serve a processor's short-term storage requirements, memory serves the processor's medium- and long-term information-storage requirements. We can classify stored information as either program or data. Program information consists of the sequence of instructions that cause the processor to carry out the desired system functionality. Data information represents the values being input, output and transformed by the program.

We can store program and data together or separately. In a *Princeton architecture*, data and program words share the same memory space. In a *Harvard architecture*, the program memory space is distinct from the data memory space. Figure 3.2 illustrates these two methods. The Princeton architecture may result in a simpler hardware connection to memory, since only one connection is necessary. A Harvard architecture, while requiring two connections, can perform instruction and data fetches simultaneously, so may result in improved performance. Most machines have a Princeton architecture. The Intel 8051 is a well-known Harvard architecture.

Memory may be read-only memory (ROM) or readable and writable memory (RAM). ROM is usually much more compact than RAM. An embedded system often uses ROM for program memory, since, unlike in desktop systems, an embedded system's program does not change. Constant data may be stored in ROM, but other data of course requires RAM.

Memory may be on-chip or off-chip. On-chip memory resides on the same IC as the processor, while off-chip memory resides on a separate IC. The processor can usually access on-chip memory much faster than off-chip memory, perhaps in just one cycle, but finite IC capacity of course implies only a limited amount of on-chip memory.




Figure 3.3: Cache memory.

✓ To reduce the time needed to access (read or write) memory, a local copy of a portion of memory may be kept in a small but especially fast memory called *cache*, as illustrated in Figure 3.3. Cache memory often resides on-chip and often uses fast but expensive static RAM technology rather than slower but cheaper dynamic RAM (see Chapter 5). Cache memory is based on the principle that if at a particular time a processor accesses a particular memory location, then the processor will likely access that location and immediate neighbors of the location in the near future. Thus, when we first access a location in memory, we copy that location and some number of its neighbors (called a *block*) into cache, and then access the copy of the location in cache. When we access another location, we first check a cache table to see if a copy of the location resides in cache. If the copy does reside in cache, we have a *cache hit*, and we can read or write that location very quickly. If the copy does not reside in cache, we have a *cache miss*, so we must copy the location's block into cache, which takes a lot of time. Thus, for a cache to be effective in improving performance, the ratio of hits to misses must be very high, requiring intelligent caching schemes. Caches are used for both program memory (often called instruction cache, or *I-cache*) as well as data memory (often called *D-cache*).

3.3 Operation

Instruction Execution

We can think of a microprocessor's execution of instructions as consisting of several basic stages:

1. *Fetch instruction*: the task of reading the next instruction from memory into the instruction register.

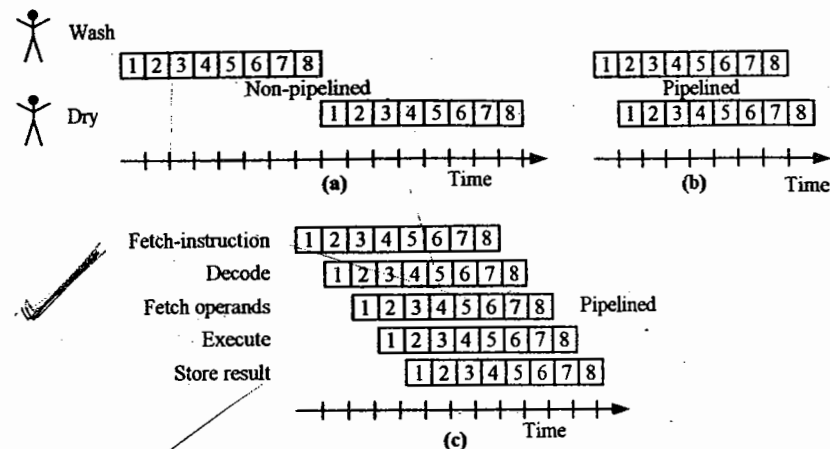


Figure 3.4: Pipelining: (a) nonpipelined dish cleaning, (b) pipelined dish cleaning, (c) pipelined instruction execution.

2. **Decode instruction:** the task of determining what operation the instruction in the instruction register represents (e.g., add, move, etc.).
3. **Fetch operands:** the task of moving the instruction's operand data into appropriate registers.
4. **Execute operation:** the task of feeding the appropriate registers through the ALU and back into an appropriate register.
5. **Store results:** the task of writing a register into memory.

If each stage takes one clock cycle, then we can see that a single instruction may take several cycles to complete.

Pipelining

Pipelining is a common way to increase the instruction throughput of a microprocessor. We first make a simple analogy of two people approaching the chore of washing and drying eight dishes. In one approach, the first person washes all eight dishes, and then the second person dries all eight dishes. Assuming 1 minute per dish per person, this approach requires 16 minutes. The approach is clearly inefficient since at any time only one person is working and the other is idle. Obviously, a better approach is for the second person to begin drying the first dish immediately after it has been washed. This approach requires only 9 minutes — 1 minute for the first dish to be washed, and then 8 more minutes until the last dish is finally dry. We refer to this latter approach as “pipelined.”

Each dish is like an instruction, and the two tasks of washing and drying are like the five stages listed earlier. By using a separate unit (each akin to a person) for each stage, we can pipeline instruction execution. After the instruction fetch unit fetches the first instruction, the

equal length stages

decode unit decodes it while the instruction fetch unit simultaneously fetches the next instruction. The idea of pipelining is illustrated in Figure 3.4. Note that for pipelining to work well, instruction execution must be decomposable into roughly equal length stages, and instructions should each require the same number of cycles.

Branches pose a problem for pipelining, since we don't know the next instruction until the current instruction has reached the execute stage. One solution is to stall the pipeline when a branch is in the pipeline, waiting for the execute stage before fetching the next instruction. An alternative is to guess which way the branch will go and fetch the corresponding instruction next; if right, we proceed with no penalty, but if we find out in the execute stage that we were wrong, we must ignore all the instructions fetched since the branch was fetched, thus incurring a penalty. Modern pipelined microprocessors often have very sophisticated branch predictors built in.

Superscalar and VLIW Architectures

We can use multiple ALUs to further speed up a processor. A *superscalar* microprocessor can execute two or more scalar operations in parallel, requiring two or more ALUs. A scalar operation transforms one or two numbers, as opposed to vector or matrix operations that transform entire sets of numbers. Some superscalar microprocessors require that the instructions be ordered statically (at compile time), while others may reorder the instructions dynamically (during runtime) to make use of the additional ALUs. A *VLIW* (very long instruction word) architecture is a type of static superscalar architecture that encodes several (perhaps four or more) operations in a single machine instruction.

3.4 Programmer's View

A programmer writes the program instructions that carry out the desired functionality on the general-purpose processor. The programmer may not actually need to know detailed information about the processor's architecture or operation, but instead may deal with an architectural abstraction, which hides much of that detail. The level of abstraction depends on the level of programming. We can distinguish between two levels of programming. The first is assembly-language programming, in which one programs in a language representing processor-specific instructions as mnemonics. The second is structured-language programming, in which one programs in a language using processor-independent instructions. A compiler automatically translates those instructions to processor-specific instructions. Ideally, the structured-language programmer would need no information about the processor architecture, but in embedded systems, the programmer must usually have at least some awareness, as we shall discuss.

Actually, we can define an even lower programming level, machine-language programming, in which the programmer writes machine instructions in binary. This level has become extremely rare due to the advent of assemblers. Machine-language-programmed computers often had rows of lights representing to the programmer the current binary instructions being executed. Today's computers look more like boxes or refrigerators, but

Instruction 1	opcode	operand1	operand2
Instruction 2	opcode	operand1	operand2
Instruction 3	opcode	operand1	operand2
Instruction 4	opcode	operand1	operand2
	...		

Figure 3.5: Instructions stored in memory.

they do not make for interesting movie props, so you may notice that in the movies, computers with rows of blinking lights live on.

Instruction Set

The assembly-language programmer must know the processor's instruction set. The instruction set describes the bit configurations allowed in the IR, indicating the atomic processor operations that the programmer may invoke. Each such configuration forms an assembly instruction, and a sequence of such instructions forms an assembly program, stored in a processor's memory, as illustrated in Figure 3.5.

An instruction typically has two parts, an opcode field and operand fields. An opcode specifies the operation to take place during the instruction. We can classify instructions into three categories. Data-transfer instructions move data between memory and registers, between input/output channels and registers, and between registers themselves. Arithmetic/logical instructions configure the ALU to carry out a particular function, move data from the registers through the ALU, and move data from the ALU back to a particular register. Branch instructions determine the address of the next program instruction, based possibly on datapath status signals.

Branches can be further categorized as being unconditional jumps, conditional jumps or procedure call and return instructions. Unconditional jumps always determine the address of the next instruction, while conditional jumps do so only if some condition evaluates to true, such as a particular register containing zero. A call instruction, in addition to indicating the address of the next instruction, saves the address of the current instruction so that a subsequent return instruction can jump back to the instruction immediately following the most recent invoked call instruction. This pair of instructions facilitates the implementation of procedure/function call semantics of high-level programming languages.

An operand field specifies the location of the actual data that takes part in an operation. Source operands serve as input to the operation, while a destination operand stores the output. The number of operands per instruction varies among processors. Even for a given processor, the number of operands per instruction may vary depending on the instruction type.

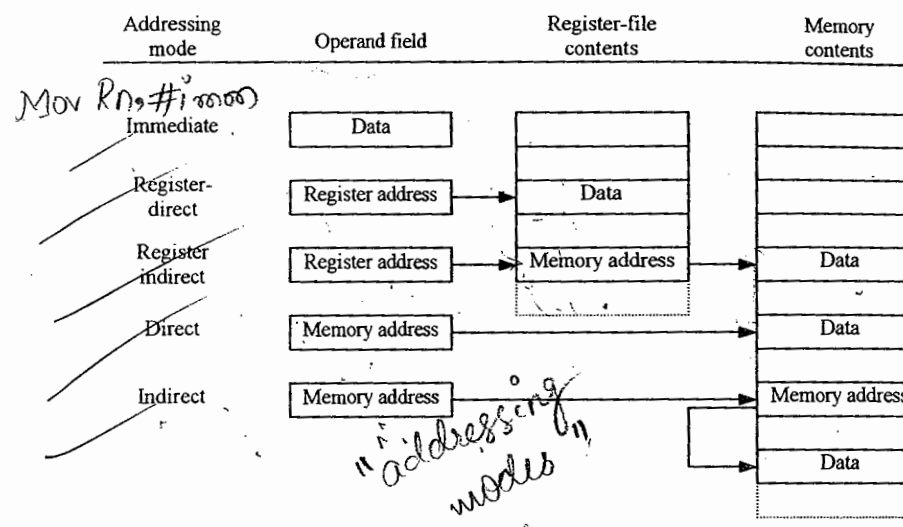


Figure 3.6: Addressing modes.

The operand field may indicate the data's location through one of several addressing modes, illustrated in Figure 3.6. In *immediate* addressing, the operand field contains the data itself. In *register* addressing, the operand field contains the address of a datapath register in which the data resides. In *register-indirect* addressing, the operand field contains the address of a register, which in turn contains the address of a memory location in which the data resides. In *direct* addressing, the operand field contains the address of a memory location in which the data resides. In *indirect* addressing, the operand field contains the address of a memory location, which in turn contains the address of a memory location in which the data resides. Those familiar with structured languages may note that direct addressing implements regular variables, and indirect addressing implements pointers. In *inherent or implicit* addressing, the particular register or memory location of the data is implicit in the opcode; for example, the data may reside in a register called the "accumulator." In *indexed* addressing, the direct or indirect operand must be added to a particular implicit register to obtain the actual operand address. Jump instructions may use *relative* addressing to reduce the number of bits needed to indicate the jump address. A relative address indicates how far to jump from the current address, rather than indicating the complete address. Such addressing is very common since most jumps are to nearby instructions.

Ideally, the structured-language programmer would not need to know the instruction set of the processor. However, nearly every embedded system requires the programmer to write at least some portion of the program in assembly language. Those portions may deal with low-level input/output operations with devices outside the processor, like a display device.

Assembly instruct.	First byte	Second byte	Operation
MOV Rn, direct	0000 Rn	direct	Rn = M(direct)
MOV direct, Rn	0001 Rn	direct	M(direct) = Rn
MOV @Rn, Rm	0010 Rn	Rm	M(Rn) = Rm
MOV Rn, #immed.	0011 Rn	immediate	Rn = immediate
ADD Rn, Rm	0100 Rn	Rm	Rn = Rn + Rm
SUB Rn, Rm	0101 Rn	Rm	Rn = Rn - Rm
JZ Rn, relative	0110 Rn	relative	PC = PC + relative (only if Rn is 0)
	opcode	operands	

Figure 3.7: A simple (trivial) instruction set.

Such a device may require specific timing sequences of signals in order to receive data, and the programmer may find that writing assembly code achieves such timing most conveniently. A *driver* routine is a portion of a program written specifically to communicate with, or drive, another device. Since drivers are often written in assembly language, the structured-language programmer may still require some familiarity with at least a subset of the instruction set.

Figure 3.7 shows a (trivial) instruction set with four data transfer instructions, two arithmetic instructions, and one branch instruction, for a hypothetical processor. Figure 3.8(a) shows a program written in C that adds the numbers 1 through 10. Figure 3.8(b) shows that same program written in assembly language using the given instruction set.

Program and Data Memory Space

The embedded systems programmer must be aware of the size of the available memory for program and for data. For example, a particular processor may have a 64K program space, and a 64K data space. The programmer must not exceed these limits. In addition, the programmer will probably want to be aware of on-chip program and data memory capacity, taking care to fit the necessary program and data in on-chip memory if possible.

Registers

Assembly-language programmers must know how many registers are available for general-purpose data storage. They must also be familiar with other registers that have special

	0	MOV R0, #0;	// total = 0	
	1	MOV R1, #10;	// i = 10	i = 10
	2	MOV R2, #1;	// constant 1	
	3	MOV R3, #0;	// constant 0	
int total = 0;				
for (int i=10; i!=0; i--)				
total += i;				
// next instructions...				
Loop:	JZ R1, Next;	// Done if i=0	i = 0	
5	ADD R0, R1;	// total += i		
6	SUB R1, R2;	// i--	i =	
7	JZ R3, Loop;	// Jump always		
Next:	// next instructions...			
(a)	(b)			

Figure 3.8: Sample programs: (a) C program, (b) equivalent assembly program.

functions. For example, a base register may exist, which permits the programmer to use a data-transfer instruction where the processor adds an operand field to the base register to obtain an actual memory address.

Other special-function registers must be known by both the assembly-language and the structured-language programmer. Such registers may be used for configuring built-in timers, counters, and serial communication devices, or for writing and reading external pins.

I/O

The programmer should be aware of the processor's input and output (I/O) facilities, with which the processor communicates with other devices. One common I/O facility is parallel I/O, in which the programmer can read or write a port (a collection of external pins) by reading or writing a special-function register. Another common I/O facility is a system bus, consisting of address and data ports that are automatically activated by certain addresses or types of instructions. I/O methods will be discussed further in Chapter 6.

Interrupts

An *interrupt* causes the processor to suspend execution of the main program and jump to an *interrupt service routine* (ISR) that fulfills a special, short-term processing need. In particular, the processor stores the current PC and sets it to the address of the ISR. After the ISR completes, the processor resumes execution of the main program by restoring the PC. The programmer should be aware of the types of interrupts supported by the processor (described in Chapter 6), and must write ISRs when necessary. The assembly-language programmer places each ISR at a specific address in program memory. The structured-language programmer must do so also; some compilers allow a programmer to force a procedure to


```

CheckPort      proc
    push ax      ; save the content
    push dx      ; save the content
    mov dx, 3BCh + 1 ; base + 1 for register #1
    in al, dx     ; read register #1
    and al, 10h   ; mask out all but bit # 4
    cmp al, 0     ; is it 0?
    jne SwitchOn ; if not, we need to turn the LED on

SwitchOff:
    mov dx, 3BCh + 0 ; base + 0 for register #0
    in al, dx         ; read the current state of the port
    and al, 0Fh       ; clear first bit (masking)
    out dx, al        ; write it out to the port
    jmp Done          ; we are done

SwitchOn:
    mov dx, 3BCh + 0 ; base + 0 for register #0
    in al, dx         ; read the current state of the port
    or al, 01h        ; set first bit (masking)
    out dx, al        ; write it out to the port

Done:  pop dx      ; restore the content
       pop ax      ; restore the content
CheckPort endp

extern "C" CheckPort(void); // defined in assembly above
void main(void) {
    while( 1 )
        CheckPort();
}

```

Figure 3.9: PC parallel port example.

start at a particular memory location, while others recognize predefined names for particular ISRs.

For example, we may need to record the occurrence of an event from a peripheral device, such as the pressing of a button. We record the event by setting a variable in memory when that event occurs, although the user's main program may not process that event until later. Rather than requiring the user to insert checks for the event throughout the main program, the programmer merely writes an interrupt service routine and associates it with an input pin connected to the button. The processor will then call the routine automatically when the button is pressed.

Example: Assembly-Language Programming of Device Drivers

This example provides an application of assembly language programming of a low-level driver, showing how the parallel port of a PC can be used to perform digital I/O. The code is

LPT Connector Pin	I/O Direction	Register Address
1	Output	0 th bit of register #2
2-9	Output	0 th -7 th bit of register #0
10, 11, 12, 13, 15	Input	6, 7, 5, 4, 3 th bit of register #1
14, 16, 17	Output	1, 2, 3 rd bit of register #2

Figure 3.10: PC parallel port signals and associated registers.

given in Figure 3.9. Writing and reading three special registers accomplishes parallel communication on the PC. Those three registers are actually in an 8255A Peripheral Interface Controller chip. In unidirectional mode, (default power-on-reset mode), this device is capable of driving 12 output and 5 input lines. In Figure 3.10, we give the parallel port (known as LPT) connector pin numbers and the corresponding register location.

A switch is connected to input pin number 13 of the parallel port. A light-emitting diode (LED) is connected to output pin number 2. Our program, running on the PC, should monitor the input switch and turn the LED on/off accordingly.

Figure 3.9 gives the code for such a program, in x86 assembly language. Note that the *in* and *out* assembly instructions read and write the internal registers of the 8255A. Both instructions take two operands, address and data. The address specifies the register we are trying to read or write. This address is calculated by adding the address of the device, called the *base address*, to the address of the particular register as given in Figure 3.9. In most PCs, the base address of LPT1 is at 3BC hex (though not always). The second operand is the data. For the *in* instruction, the content of this 8-bit operand will be written to the addressed register. For the *out* instruction, the content of the addressed 8-bit register will be read into this operand.

The program makes use of masking, something quite common during low-level I/O. A *mask* is a bit-pattern designed such that ANDing it with a data item D yields a specific part of D. For example, a mask of 00001111 can be used to yield bits 3 through 0 (e.g., 00001111 AND 10101010 yields 00001010). A mask of 00010000, or 10h in hexadecimal format, would yield bit 4.

In Figure 3.9, we have broken our program in two source files, assembly and C. The assembly program implements the low-level I/O to the parallel port and the C program implements the high-level application.

Operating System

An *operating system* is a layer of software that provides low-level services to the application layer, a set of one or more programs executing on the CPU consuming and producing input and output data. The task of managing the application layer involves the loading and executing of programs, sharing and allocating system resources to these programs, and

```

DB file_name "out.txt" - store file name

MOV R0, 1324           - system call "open" id
MOV R1, file_name      - address of file-name
INT 34                 - cause a system call
JZ R0, L1              - if zero -> error

... read the file
JMP L2                 - bypass error condition
L1:
... handle the error

L2:

```

Figure 3.11: System call invocation.

protecting these allocated resources from corruption by non-owner programs. One of the most important resource of a system is the central processing unit (CPU), which is typically shared among a number of executing programs. The operating system is responsible for deciding what program is to run next on the CPU and for how long. This is called *process* (or *task*) *scheduling* and it is determined by the operating system's preemption policy. Another very important resource is memory, including disk storage, which is also shared among the applications running on the CPU.

In addition to implementing an environment for management of high-level application programs, the operating system provides the software required for servicing various hardware-interrupts, and provides device drivers for driving the peripheral devices present in the system. Typically, on startup, an operating system initializes all peripheral devices, such as disk controllers, timers, and input/output devices and installs hardware interrupt service routines (ISRs) to handle various signals generated by these devices. Then it installs software interrupts (interrupts generated by the software) to process system calls (calls made by high-level applications to request operating system services) as described next.

A *system call* is a mechanism for an application to invoke the operating system. It is analogous to a procedure or function call, as in high-level programming languages. When a program requires some service from the operating system, it generates a predefined software interrupt that is serviced by the operating system. Parameters specific to the requested services are typically passed from (to) the application program to (from) the operating system through CPU registers. Figure 3.11 illustrates how the file "open" system call may be invoked, in assembly, by a program. Languages like C and Pascal provide wrapper functions around the system-calls to provide a high-level mechanism for performing system calls.

In summary, the operating system abstracts away the details of the underlying hardware and provides the application layer an interface to the hardware through the system call mechanism.

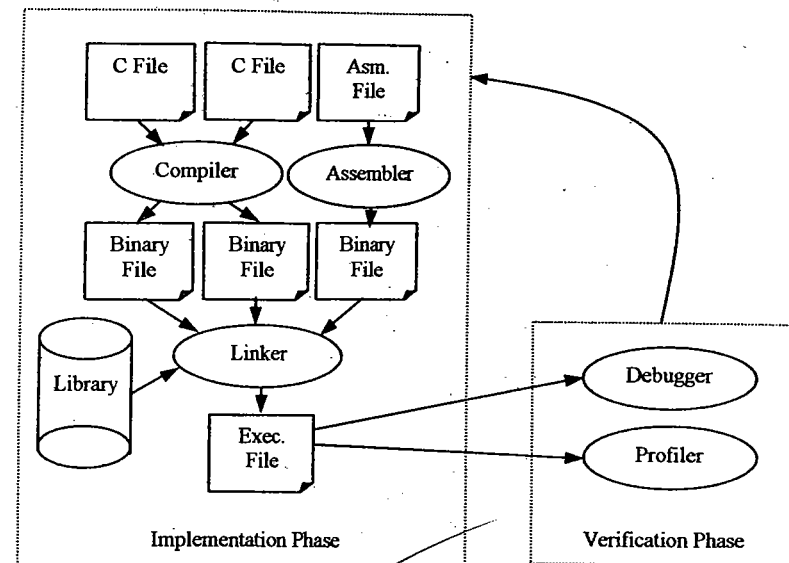


Figure 3.12: Software development process.

3.5 Development Environment

In this section, we take a look at the general software design tools that are used by embedded system designers in design, test, and debugging of embedded software.

Design Flow and Tools

Several software and hardware tools commonly support the programming of general-purpose processors. First, we must distinguish between two processors we deal with when developing an embedded system. One processor is the *development processor*, on which we write and debug our program. This processor is part of our desktop computer. The other processor is the *target processor*, to which we will send our program and which will form part of our embedded system's implementation. For example, we may develop our system on a Pentium processor but use a Motorola 68HC11 as our target processor. Of course, sometimes the two processors happen to be the same, but this is mostly a coincidence.

Programming of an embedded system's processor is similar to writing a program that runs on your desktop computer, with some subtle but important differences. Figure 3.12 depicts the standard software development process. The general design flow for programming

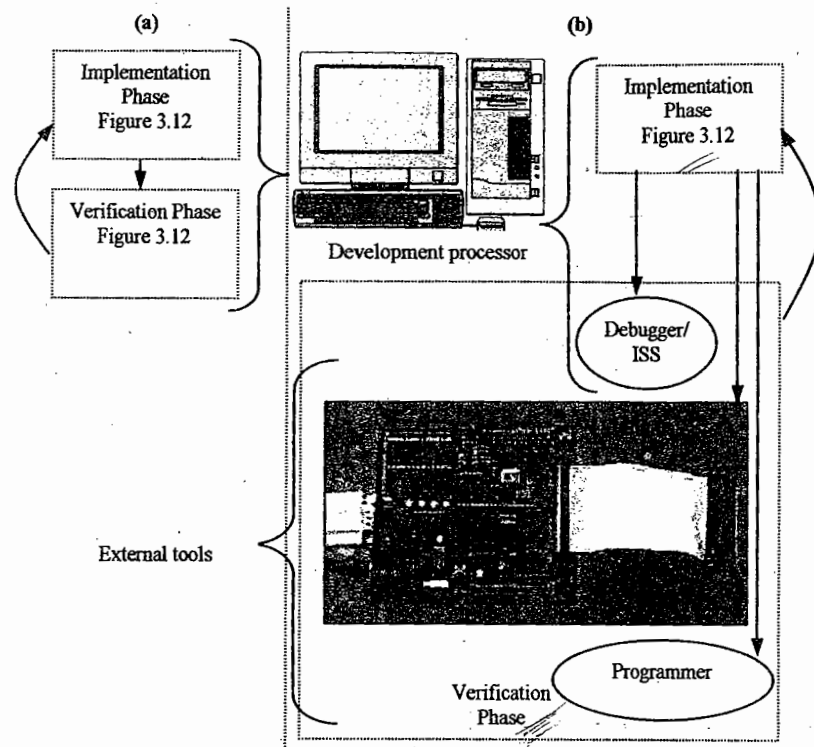


Figure 3.13: Software design process: (a) desktop, (b) embedded.

applications that run on a desktop computer starts with writing our source code, possibly organized in a number of files for modularity, using an editor. Then we compile or assemble the code in each file, using a compiler or assembler, into corresponding binary files. Next, using a linker, we combine these binary files into a final executable. These tasks, collectively, can be considered as the implementation phase. Next, we test our program by running the executable file under the command of a debugger. Sometimes we may use a profiler to pinpoint the performance bottlenecks of your program. During this phase, if we discover errors or performance bottlenecks, we return to the implementation phase, make improvements, and repeat the process.

Typically, all of these tools have been combined into a single *integrated development environment* (IDE), which greatly simplifies the design process. Figure 3.13(b) shows the design flow for embedded software development, in contrast to the design flow for desktop applications in Figure 3.13(a). Here, the implementation phase, which is the process of

editing, compiling, assembling, and linking our program, is the same as that used for designing applications intended for our desktop computer. We perform these on our development computer using cross-compilers, cross-assemblers, etc. It is the verification phase (i.e., the process of testing the final executable) that is greatly different in embedded systems. In the following paragraphs, we will describe each of the development tools in greater detail.

(a) Assemblers translate assembly instructions to binary machine instructions. In addition to just replacing opcode and operand mnemonics by binary equivalents, an assembler may also translate symbolic labels into actual addresses. For example, a programmer may add a symbolic label *END* to an instruction *A* and may reference *END* in a branch instruction. The assembler determines the actual binary address of *A* and replaces references to *END* by this address. The mapping of assembly instructions to machine instructions is one-to-one.

(b) Compilers translate structured programs into machine (or assembly) programs. Structured programming languages possess high-level constructs that greatly simplify programming, such as loop constructs, so each high-level construct may translate to several or tens of machine instructions. Compiler technology has advanced tremendously over the past decades, applying numerous program optimizations, often yielding very size and performance efficient code. A cross compiler executes on one processor (our development processor) but generates code for a different processor (our target processor). Cross compilers are extremely common in embedded system development.

(c) Linkers allow a programmer to create a program in separately assembled or compiled files; it combines the machine instructions of each into a single program, perhaps incorporating instructions from standard library routines. A linker designed for embedded processors will also try to eliminate binary code associated with uncalled procedures and functions as well as memory allocated to unused variables in order to reduce the overall program footprint.

Example: Instruction-Set Simulator for a Simple Processor

An *instruction-set simulator* is a program that runs on one processor and executes the instructions of another processor. In this example, we design an instruction-set simulator for the simple processor of Figure 3.7. Our program takes as input a file name containing binary instructions of our simple processor. The code for this instruction-set simulator is given in Figure 3.14.

Testing and Debugging

Generally, the testing and debugging phase of developing programs is a major part of the overall design processes. This is especially true when the program is being developed to run in an embedded system. For example, it is not acceptable for your car's engine management system to require occasional rebooting because of a software hang up. Programming is an error prone activity, and it is inevitable that there will exist errors and bugs in writing any reasonably large program. The most common method of verifying the correctness of a program is running (executing) it with ample input data that check the program's behavior.

```

#include <stdio.h>
typedef struct {
    unsigned char first_byte, second_byte;
} instruction;
instruction program[1024]; // this is our instruction memory
unsigned char memory[256]; // this is our data memory
void run_program(int num_bytes) {
    int pc = -1;
    unsigned char reg[16], fb, sb;
    while( ++pc < (num_bytes / 2) ) {
        fb = program[pc].first_byte;
        sb = program[pc].second_byte;
        switch( fb >> 4 ) {
            case 0: reg[fb & 0x0f] = memory[sb]; break;
            case 1: memory[sb] = reg[fb & 0x0f]; break;
            case 2: memory[reg[fb & 0x0f]] = reg[sb >> 4]; break;
            case 3: reg[fb & 0x0f] = sb; break;
            case 4: reg[fb & 0x0f] += reg[sb >> 4]; break;
            case 5: reg[fb & 0x0f] -= reg[sb >> 4]; break;
            case 6: if(reg[fb & 0x0f] == 0) pc += sb; break;
            default: return -1;
        }
    }
    return 0;
}

int main(int argc, char *argv[]) {
    FILE* ifs;
    if( argc != 2 || (ifs = fopen(argv[1], "rb") == NULL) ) return -1;
    if (run_program(fread(program, sizeof(program) = 0) {
        print_memory_contents();
        return(0);
    }
    else return(-1);
}

```

Figure 3.14: Instruction-set simulator implementation.

especially using boundary cases. This is relatively easy to do when developing programs that run on your desktop computer.

For embedded system programmers, this task is a little more challenging. Specifically, a program running in an embedded system most often needs to be real-time. For example, our engine management program must generate pulses that actuate the fuel-injectors with a timely and calculated pattern. A distinguishing characteristic of a real-time system is that it must compute correct results within a predetermined amount of time, while a non-real-time system only needs to compute correct results. In addition, a program running in an embedded system works in conjunction with many other components of that system as well as interacts with the

environment where the embedded system is to function. Hence, debugging a program running in an embedded system requires having control over time, as well as control over the environment and the ability to trace or follow the execution of the program, in order to detect errors. In the remaining paragraphs, we take a look at some tools and methods to help us do just that. These tools, for the most part, enable us to execute and observe the behavior of our programs.

Debuggers help programmers evaluate and correct their programs. They run on the development processor and support stepwise program execution) executing one instruction and then stopping, proceeding to the next instruction when instructed by the user. They permit execution up to user-specified breakpoints, which are instructions that when encountered cause the program to stop executing. Whenever the program stops, the user can examine values of various memory and register locations. A source-level debugger enables step-by-step execution in the source program language, whether assembly language or a structured language. A good debugging capability is crucial, as today's programs can be quite complex and hard to write correctly. Since debuggers are programs that run on your development processor but execute code designed for your target processor, they always mimic or simulate the function of the target processor. These debuggers are also known as instruction-set simulators (ISS) or virtual machines (VM).

Emulators support debugging of the program while it executes on the target processor. An emulator typically consists of a debugger coupled with a board connected to the desktop processor via a cable. The board consists of the target processor plus some support circuitry (often another processor). The board may have another cable with a device having the same pin configuration as the target processor, allowing one to plug the device into a real embedded system. Such an in-circuit emulator enables one to control and monitor the program's execution in the actual embedded system circuit. In-circuit emulators are available for nearly any processor intended for embedded use, although they can be quite expensive if they are to run at real speeds.

Device programmers download a binary machine program from the development processor's memory into the target processor's memory. Once the target processor has been programmed, the entire embedded system can be tested in its most realistic form (i.e., it can be executed in its environment and the behavior observed in a realistic way). For example, a car equipped with our engine management system can be taken out for a drive!

Revisiting Figure 3.12, we see that programs intended for embedded systems can be tested in three ways, namely, debugging using an ISS, emulation using an emulator, and field testing by downloading the program directly into the target processor. The difference between these three methods is as follows. The design cycle using a debugger based on an ISS running on the development computer is fast, but it is inaccurate since it can only interact with the rest of the system and the environment to a limited degree. The design cycle using an emulator is a little longer, since the code must be downloaded into the emulator hardware; however, the emulator hardware can interact with the rest of the system, hence can allow for more accurate testing. The design cycle using a programmer to download the program into the target processor is the longest of all. Here, the target processor must be removed from its system and put into the programmer, programmed, and returned to the system. However, this method will

enable the system to interact with its environment more freely, hence provides the highest execution accuracy but little debug control.

The availability of low-cost or high-quality development environments for a processor often heavily influences the choice of a processor.

3.6 Application-Specific Instruction-Set Processors (ASIPs)

Today's embedded applications, such as high definition TV, require high computing power and very specific functionality. The performance, power, cost, or size demands of these applications cannot always be dealt with efficiently by using general-purpose processors. Nonetheless, the inflexibility of custom single-purpose processors is often too prohibitive. A solution is to use an instruction-set processor that is specific to that application or application domain. Because these ASIPs are instruction-set processors, they can be programmed by writing software, resulting in short time-to-market and good flexibility, while the performance and other constraints may be efficiently satisfied.

As with most other aspects of embedded systems design, there is a trade-off here. Instruction-set processors and the associated software tools (compilers, linkers, etc.) are very expensive to develop; therefore, they are expensive to integrate into low-cost embedded systems. In contrast, the large applicability and resulting cost amortization of general-purpose processors make them very cost effective solutions in most embedded systems. ASIPs tend to come in three major varieties, namely, microcontrollers, which are specific to applications that perform a large amount of control-oriented tasks, digital signal processors (DSPs), which are specific to applications that process large amounts of data, and everything else, which are less general ASIPs.

Microcontrollers

Numerous processor IC manufacturers market devices specifically for the control-dominated embedded systems domain. These devices may include several features. First, they may include several peripheral devices, such as timers, analog-to-digital converters, and serial communication devices, on the same IC as the processor. Second, they may include some program and data memory on the same IC. Third, they may provide the programmer with direct access to a number of pins of the IC. Fourth, they may provide specialized instructions for common embedded system control operations, such as bit-manipulation operations. A *microcontroller* is a device possessing some or all of these features.

Incorporating peripherals and memory onto the same IC reduces the number of required ICs, resulting in compact and low-power implementations. Providing pin access allows programs to easily monitor sensors, to set actuators, and to transfer data with other devices. Providing specialized instructions improves performance for embedded systems applications. Thus, microcontrollers can be considered ASIPs to some degree.

Many manufacturers market devices referred to as "embedded processors." The difference between embedded processors and microcontrollers is not clear, although we note that the former term seems to be used more for large (32-bit) processors.

Digital Signal Processors (DSP)

Digital signal processors (DSPs) are processors that are highly optimized for processing large amounts of data. The source of this large amount of data is some form of digitized signal, like a photo image captured by a digital camera, a voice packet going through a network router, or an audio clip played by a digital keyboard. A DSP may contain numerous register files, memory blocks, multipliers, and other arithmetic units. In addition, DSPs often provide instructions that are central to digital signal processing, such as filtering and transforming vectors or metrics of data. In a DSP, frequently used arithmetic functions, such as multiply-and-accumulate, are implemented in hardware and thus execute orders of magnitude faster than a software implementation running on a general-purpose processor. In addition, DSPs may allow for execution of some functions in parallel, resulting in a boost in performance.

As with microcontrollers, DSPs also tend to incorporate many peripherals that are useful in signal processing on a single IC. As an example, a DSP device may contain a number of analog-to-digital and digital-to-analog converters, pulse-width-modulators, direct-memory-access controllers, timers, and counters.

Many companies offer a variety of commonly used DSPs that are well supported in terms of compiler and other development tools, making them easy and cheap to integrate into most embedded systems.

Less-General ASIP Environments

In contrast to microcontrollers and DSPs, which can be used in a variety of embedded systems, IC manufacturers have designed ASIPs that are less general in nature. These ASIPs are designed to perform some very domain specific processing while allowing some degree of programmability. For example, an ASIP designed for networking hardware may be designed to be programmable with different network routing, checksum, and packet processing protocols.

3.7 Selecting a Microprocessor

The embedded system designer must select a microprocessor for use in an embedded system. The choice of a processor depends on technical and nontechnical aspects. From a technical perspective, one must choose a processor that can achieve the desired speed within certain power, size and cost constraints. Nontechnical aspects may include prior expertise with a processor and its development environment, special licensing arrangements, and so on.

Speed is a particularly difficult processor aspect to measure and compare. We could compare processor clock speeds, but the number of instructions per clock cycle may differ greatly among processors. We could instead compare instructions per second, but the complexity of each instruction may also differ greatly among processors. For example, one processor may require 100 instructions while another processor may require 300 instructions to perform the same computation.

Processor	Clock Speed	Peripherals	Bus Width	MIPS	Power	Transistors	Price
General Purpose Processors							
Intel PIII	1GHz	2x16 K L1, 256K L2, MMX	32	~900	97 W	~7 M	\$900
IBM PowerPC 750X	550 MHz	2x32K L1, 256K L2	32/64	~1300	5 W	~7 M	\$900
MIPS R5000	250 MHz	2x32K, 2 way set assoc.	32/64	NA	NA	3.6 M	NA
StrongARM SA-110	233 MHz	None	32	268	1 W	2.1 M	NA
Microcontrollers							
Intel 8051	12 MHz	4K ROM, 128 RAM, 32 I/O, Timer, UART	8	~1	~0.2 W	~10 K	\$7
Motorola 68HC811	3 MHz	4K ROM, 192 RAM, 32 I/O, Timer, WDT, SPI	8	~5	~0.1 W	~10 K	\$5
Digital Signal Processors							
TI C5416	160 MHz	128K SRAM, 3 T1 Ports, DMA, 13 ADC, 9 DAC	16/32	~600	NA	NA	\$34
Lucent DSP32C	80 MHz	16K Inst., 2K Data, Serial Ports, DMA	32	40	NA	NA	\$75

Sources: Intel, Motorola, MIPS, ARM, TI, and IBM Websites/Datasheets; Embedded Systems Programming, Nov. 1998.

Figure 3.15: General-purpose processors.

One attempt to provide a means for a fairer comparison is the *Dhrystone benchmark*. A benchmark is a program intended to be run on different processors to compare their performance. The Dhrystone benchmark was originally developed in 1984 by Reinhold Weicker specifically as a performance benchmark; it performs no useful work. It focuses on exercising a processor's integer arithmetic and string-handling capabilities. Its current version is written in C and is in the public domain. Because most processors can execute it in milliseconds, it is typically executed thousands of times, and thus a processor is said to be able to execute so many Dhrystones per second.

Another commonly used speed comparison unit, which happens to be based on the Dhrystone, is MIPS. One might think that MIPS simply means millions of instructions per second, but actually the common use of the term is based on a somewhat more complex notion. Specifically, its origin is based on the speed of Digital's VAX 11/780, thought to be the first computer able to execute one million instructions per second. A VAX 11/780 could execute 1,757 Dhrystones/second. Thus, for a VAX 11/780, 1 MIPS = 1,757 Dhrystones/second. This unit for MIPS is the one commonly used today, and it is sometimes referred to as *Dhrystone MIPS*. So if a machine today is said to run at 750 MIPS, that actually means it can execute $750 \times 1,757 = 1,317,750$ Dhrystones/second.

The use and validity of benchmark data is a subject of great controversy. There is also a clear need for benchmarks that measure performance of embedded processors. An effort underway in this area is EEMBC (pronounced "embassy"), the EDN Embedded Benchmark Consortium. The EEMBC has five benchmarking suites of programs corresponding to different embedded applications: automotive/industrial, consumer electronics, networking, office automation, and telecommunications. Each suite consists of several common algorithms found in the suite's application area. For example, two of the programs in the consumer electronics suite are JPEG compression and decompression (JPEG is a standard for still digital image compression). Another program in that suite involves infrared signal transmission and reception.

Numerous general-purpose processors have evolved in the recent years and are in common use today. In Figure 3.15, we summarize some of the features of several popular processors.

3.8 General-Purpose Processor Design

A general-purpose processor is really just a single-purpose processor whose purpose is to process instructions stored in a program memory. Therefore, we can design a general-purpose processor using the single-purpose processor design technique described in Chapter 2. While real microprocessors intended for mass production are more commonly designed using custom methods rather than the general technique of this section, using the general technique here may prove a useful exercise that will illustrate the basic unity between single-purpose and general-purpose processors.

Suppose we want to design a general-purpose processor having the basic architecture of Figure 3.1 and supporting the instruction set of Figure 3.7. We can begin by creating the FSMD shown in Figure 3.16(a) which describes the desired processor's behavior. The FSMD declares several variables for storage: a 16-bit program counter *PC*, a 16-bit instruction register *IR*, a $64K \times 16$ bit memory *M*, and a 16×16 bit register file *RF*. The FSMD's initial state, *Reset*, clears *PC* to 0. The *Fetch* state reads $M[PC]$ into *IR*. The *Decode* state does nothing but adds the extra cycle necessary for *IR* to get updated so we can then read it on an arc. Each arc leaving the *Decode* state detects a particular instruction opcode, causing a transition to the corresponding execute state for that opcode. Each execute state, like *Movl*,

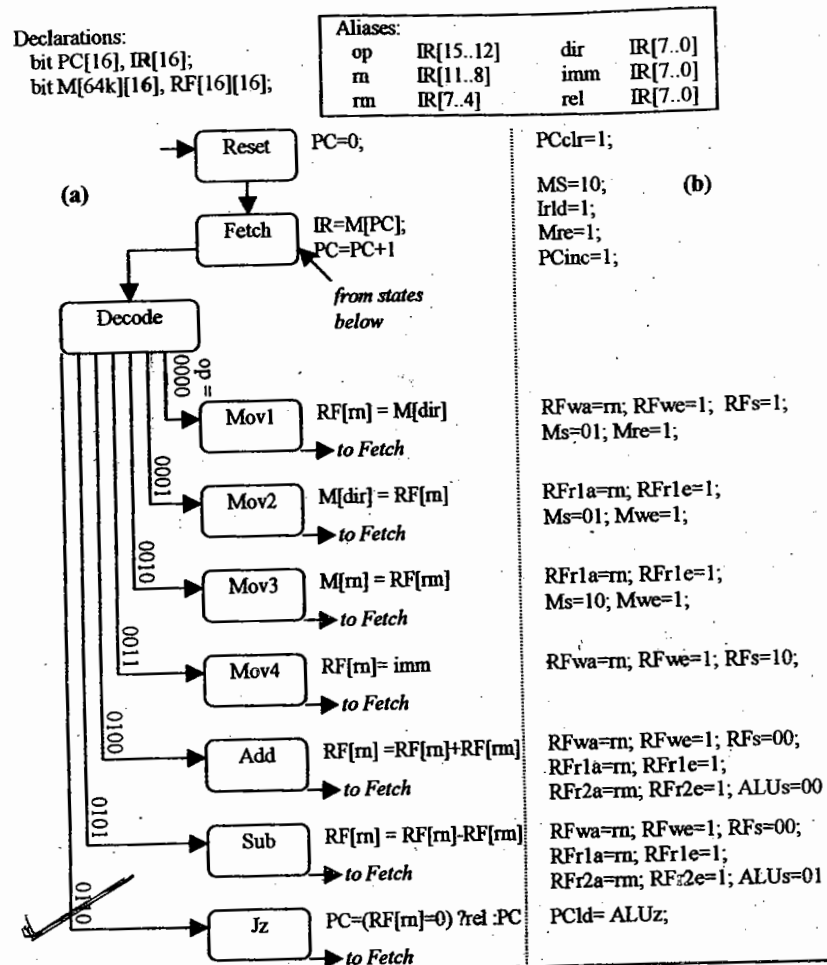


Figure 3.16: A simple microprocessor: (a) FSMD, (b) FSM operations that replace the FSMD operations after we create the datapath of Figure 3.17.

Add, and Jz, carries out the actual instruction operations by moving data between storage devices, modifying data, or updating PC.

We can now build a datapath that can carry out the operation of this FSMD, as described in Chapter 2. The datapath we create using the following steps is shown in Figure 3.17. The first step is to instantiate a storage device for each declared variable, so we instantiate

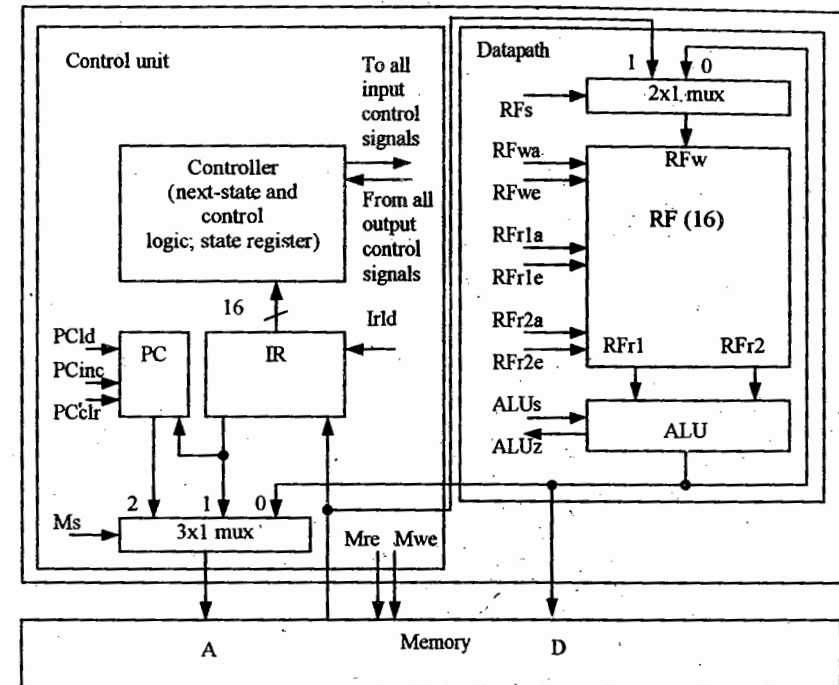


Figure 3.17: Architecture of a simple microprocessor.

registers PC and IR, memory M, and register file RF. The second step is to instantiate functional units to carry out the FSMD operations. We'll use a single ALU capable of carrying out all the operations. The third step is to add the connections among the components' ports as required by the FSMD operations, adding multiplexors when there is more than one connection being input to a port. Finally, we create unique identifiers for every control signal.

Given this datapath, we can now rewrite the FSMD as an FSM representing the datapath's controller. Each FSMD operation must be replaced by binary operations on control signals, as shown in Figure 3.16(b). The states and arcs are identical for the FSMD and FSM, and only the operations change, so we do not redraw the states and arcs in the figure. As an example of operation replacement, we replace the assignment PC = 0 in state-Reset by the control signal setting PCclr = 1.

We can use the FSM design technique of Chapter 2 to design a controller, consisting of a state register and next-state/control logic. We omit this step here.

Having just designed a simple general-purpose processor using the same technique we used to design a single-purpose processor, we can see the similarity between the two processor types. The key difference is that a single-purpose processor puts the "program" inside of its control logic, whereas a general-purpose processor keeps it in an external memory. So the program of a single-purpose processor cannot be changed once the processor has been implemented. But nevertheless, both processor types process programs. A second difference is that we design the datapath in a general-purpose processor without knowledge of what program will be put in the memory, whereas we know this program in a single-purpose processor. So the datapath of a single-purpose processor can be optimized to the program. We see that single-purpose and general-purpose processors both implement programs. Though they may differ in terms of design metrics like flexibility, power, performance, and cost, they fundamentally do the same thing.

3.9 Summary

General-purpose processors are popular in embedded systems due to several features, including low unit cost, good performance, and low NRE cost. A general-purpose processor consists of a controller and datapath, with a memory to store program and data. To use a general-purpose processor, the embedded system designer must write a program. The designer may write some parts of this program, such as driver routines, using assembly language, while writing other parts in a structured language. Thus, the designer should be aware of several aspects of the processor being used, such as the instruction set, available memory, registers, I/O facilities, and interrupt facilities. Many tools exist to support the designer, including assemblers, compilers, debuggers, device programmers, and emulators. The designer often makes use of microcontrollers, which are processors specifically targeted to embedded systems. These processors may include on-chip peripheral devices and memory, additional I/O ports, and instructions supporting common embedded system operations. The designer has a variety of processors from which to choose.

3.10 References and Further Reading

- Philips semiconductors, *80C51-based 8-bit Microcontrollers Databook*, Philips Electronics North America, 1994. Provides an overview of the 8051 architecture and on-chip peripherals, describes a large number of derivatives each with various features, describes the I2C and CAN bus protocols, and highlights development support tools.
- Rafiquzzaman, Mohamed. *Microprocessors and Microcomputer-Based System Design*. Boca Raton: CRC Press, 1995. Provides an overview of general-purpose processor architecture, along with detailed descriptions of various Intel 80xx and Motorola 68000 series processors.
- *Embedded Systems Programming*, Miller Freeman Inc., San Francisco, 1999. A monthly publication covering trends in various aspects of general-purpose processors for

embedded systems, including programming, compilers, operating systems, emulators, device programmers, microcontrollers, PLDs, and memories. An annual buyer's guide provides tables of vendors for these items, including 8/16/32/64-bit microcontrollers/microprocessors and their features.

- *Microprocessor Report*, MicroDesign Resources, California, 1999. A monthly report providing in-depth coverage of trends, announcements, and technical details, for desktop, mobile, and embedded microprocessors.
- www.eembc.org. The Web site for the EEMBC benchmark consortium.
- SIGPLAN Notices 23,8 (Aug. 1988), 49-62. Provides source for the Dhrystone benchmark version 2. Online source can be found at ftp.nosc.mil/pub/aburto.

3.11 Exercises

- 3.1 Describe why a general-purpose processor could cost less than a single-purpose processor you design yourself.
- 3.2 Detail the stages of executing the MOV instructions of Figure 3.7, assuming an 8-bit processor and a 16-bit IR and program memory following the model of Figure 3.1. For example, the stages for the ADD instruction are (1) fetch M[PC] into IR, (2) read Rn and Rm from register file through ALU configured for ADD, storing results back in Rn.
- 3.3 Add one instruction to the instruction set of Figure 3.7 that would reduce the size of our summing assembly program by 1 instruction. *Hint*: add a new branch instruction. Show the reduced program.
- 3.4 Create a table listing the address spaces for the following address sizes: (a) 8-bit, (b) 16-bit, (c) 24-bit, (d) 32-bit, (e) 64-bit.
- 3.5 Illustrate how program and data memory fetches can be overlapped in a Harvard architecture.
- 3.6 Read the entire problem before beginning. (a) Write a C program that clears an array "short int M[256]." In other words, the program sets every location to 0. *Hint*: your program should only be a couple lines long. (b) Assuming M starts at location 256 (and thus ends at location 511), write the same program in assembly language using the earlier instruction set. (c) Measure the time it takes you to perform parts a and b, and report those times.
- 3.7 Acquire a databook for a microcontroller. List the features of the basic version of that microcontroller, including key characteristics of the instruction set (number of instructions of each type, length per instruction, etc.), memory architecture and available memory, general-purpose registers, special-function registers, I/O facilities, interrupt facilities, and other salient features.
- 3.8 For the microcontroller in the previous exercise, create a table listing five existing variations of that microcontroller, stressing the features that differ from the basic version.